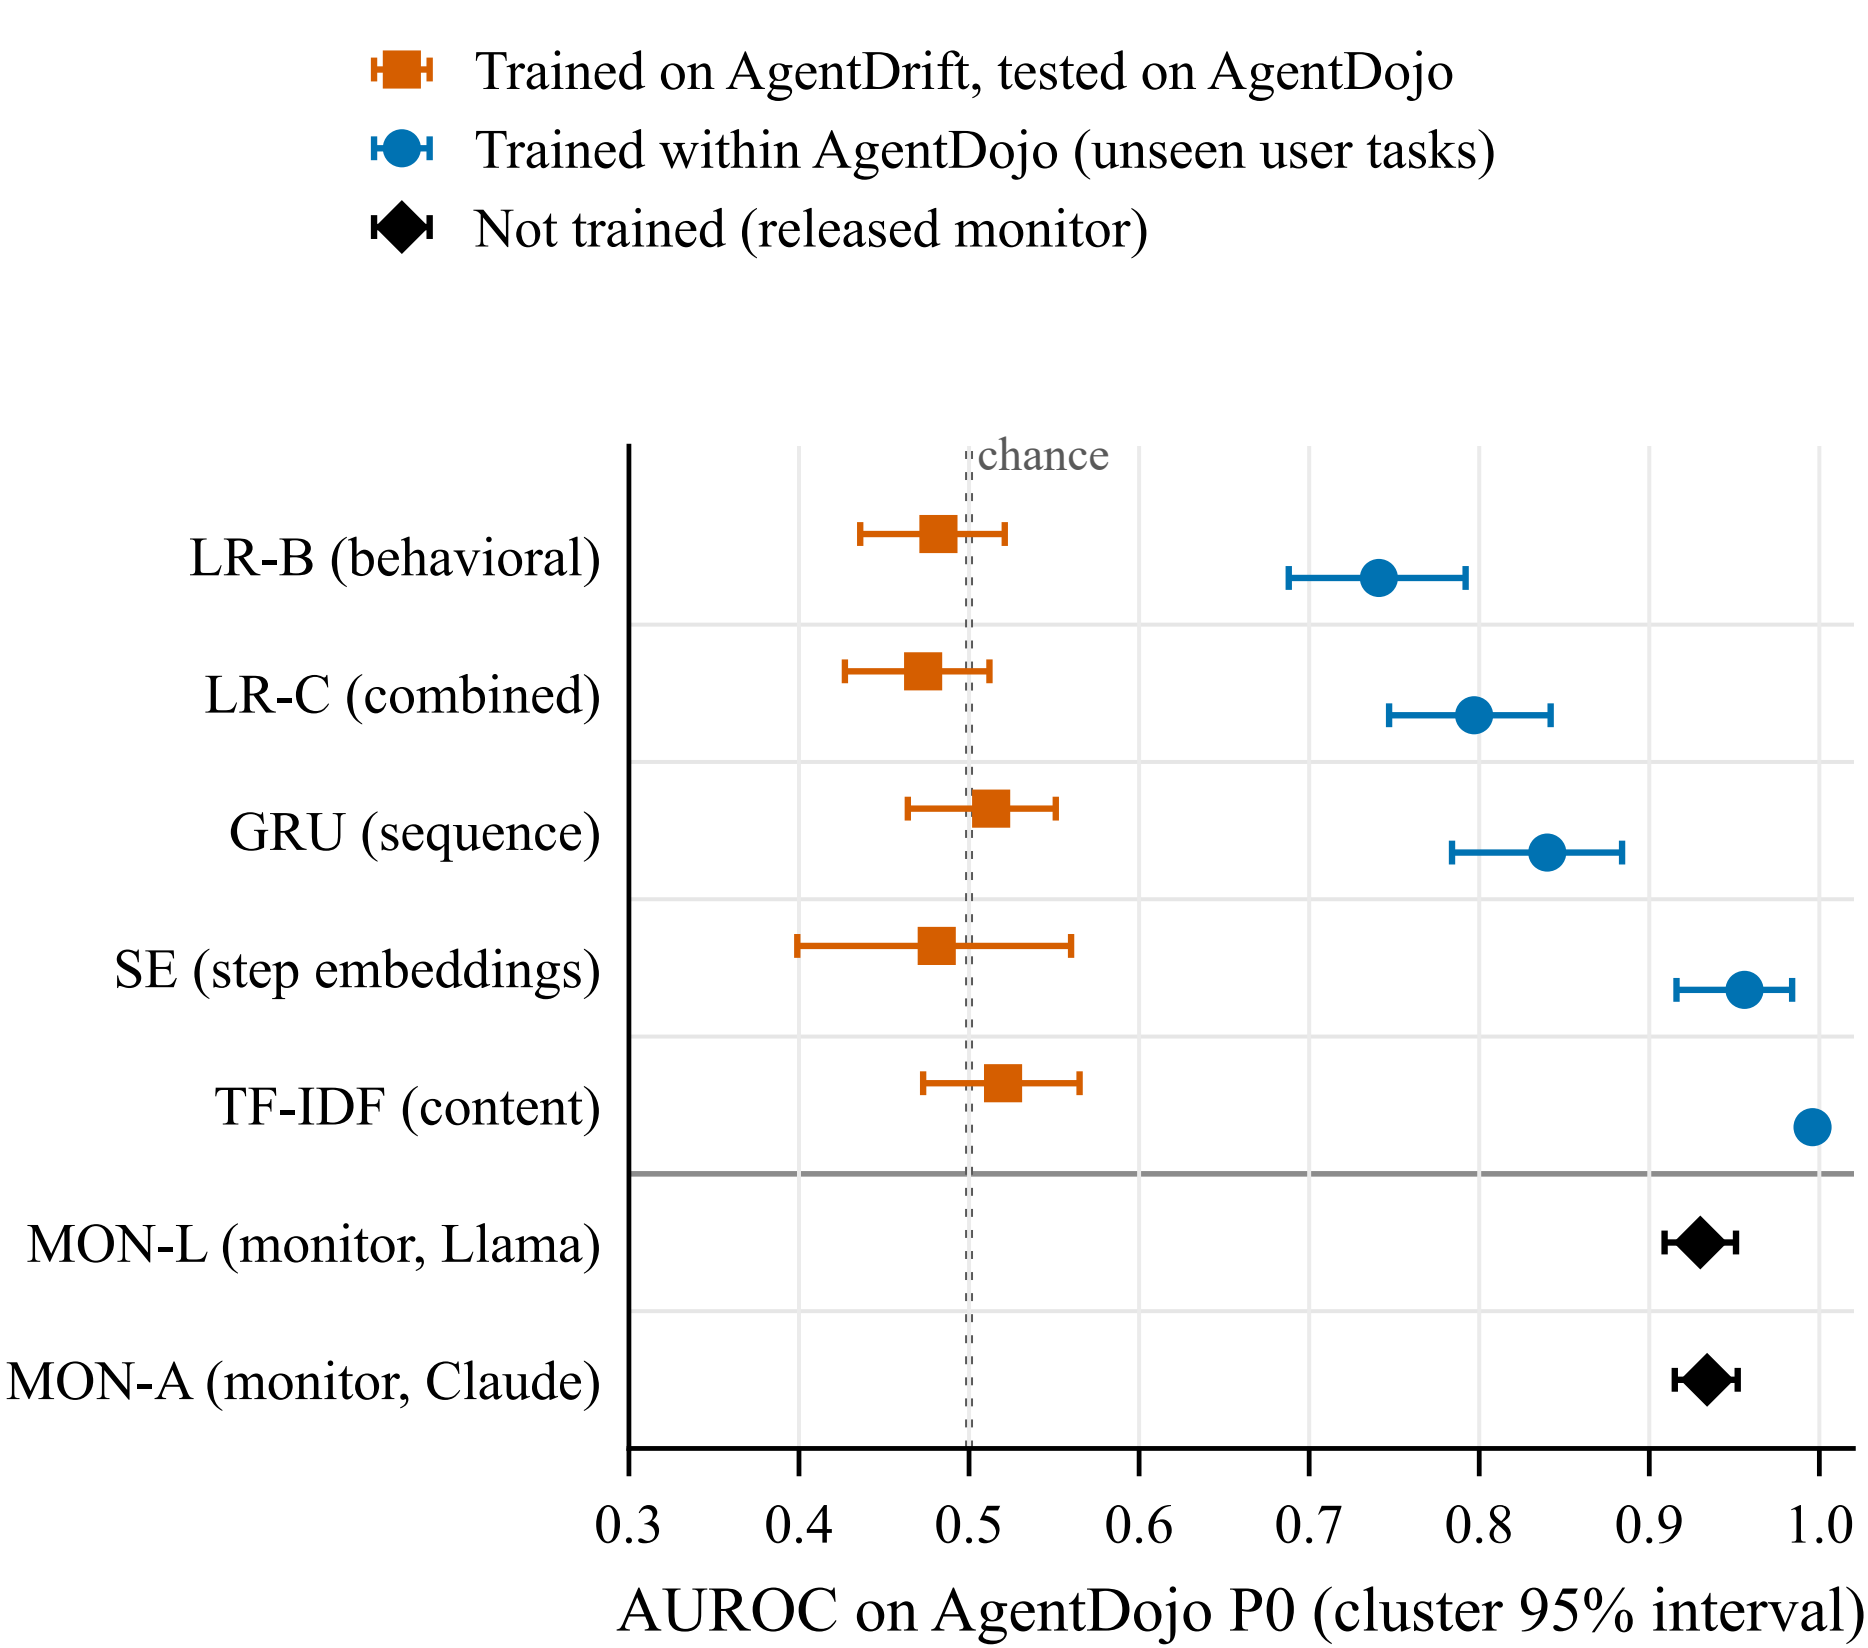
Trained on AgentDrift, tested on AgentDojo
Trained within AgentDojo (unseen user tasks)
Not trained (released monitor)
chance
LR-B (behavioral)
LR-C (combined)
GRU (sequence)
SE (step embeddings)
TF-IDF (content)
MON-L (monitor, Llama)
MON-A (monitor, Claude)
0.3
0.4
0.5
0.6
0.7
0.8
0.9
1.0
AUROC on AgentDojo P0 (cluster 95% interval)